

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

IMPORTING THE LIBRARIES :-

```
In [2]: import pandas as pd
import numpy as np
import pickle
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import sklearn
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import RandomizedSearchCV
import imblearn
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import tensorflow
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
import warnings
warnings.filterwarnings('ignore')
```

READING THE DATASET :-

```
In [3]: df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/Saran/train_u6lu')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

```
In [5]: df.drop('Loan_ID',axis=1,inplace=True)
```

```
In [6]: df.head()
```

Out[6]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantInc
--	--------	---------	------------	-----------	---------------	-----------------	----------------

0	Male	No	0	Graduate	No	5849	
1	Male	Yes	1	Graduate	No	4583	15
2	Male	Yes	0	Graduate	Yes	3000	
3	Male	Yes	0	Not Graduate	No	2583	23
4	Male	No	0	Graduate	No	6000	

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                601 non-null    object
1   Married               611 non-null    object
2   Dependents            599 non-null    object
3   Education             614 non-null    object
4   Self_Employed        582 non-null    object
5   ApplicantIncome       614 non-null    int64
6   CoapplicantIncome     614 non-null    float64
7   LoanAmount            592 non-null    float64
8   Loan_Amount_Term      600 non-null    float64
9   Credit_History        564 non-null    float64
10  Property_Area         614 non-null    object
11  Loan_Status           614 non-null    object
dtypes: float64(4), int64(1), object(7)
memory usage: 57.7+ KB
```

In [8]: `df.isnull().sum()`

Out[8]:

Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

dtype: int64

In [9]:

```
cat_cols = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', '
num_cols = ['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amou

for col in cat_cols:
    df[col] = df[col].fillna(df[col].mode()[0])

for col in num_cols:
    df[col] = df[col].fillna(df[col].mean())
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: Gender                0
Married                0
Dependents             0
Education              0
Self_Employed         0
ApplicantIncome        0
CoapplicantIncome      0
LoanAmount             0
Loan_Amount_Term       0
Credit_History         0
Property_Area          0
Loan_Status            0
dtype: int64
```

```
In [11]: for col in cat_cols:
          print(col)
          print(df[col].unique())
          print('\n')
```

```
Gender
['Male' 'Female']
```

```
Married
['No' 'Yes']
```

```
Dependents
['0' '1' '2' '3+']
```

```
Education
['Graduate' 'Not Graduate']
```

```
Self_Employed
['No' 'Yes']
```

```
Property_Area
['Urban' 'Rural' 'Semiurban']
```

```
Loan_Status
['Y' 'N']
```

HANDLING CATEGORICAL VALUES :-

```
In [12]: df.Gender.replace(['Male', 'Female'], [0, 1], inplace=True)
```

```
In [13]: df.Married.replace(['No', 'Yes'], [0, 1], inplace=True)
```

```
In [14]: df.Dependents = df.Dependents.str.replace('+', '')
```

```
In [15]: df.Education.replace(['Graduate', 'Not Graduate'],[0,1],inplace=True)
```

```
In [16]: df.Self_Employed.replace(['No', 'Yes'],[0,1],inplace=True)
```

```
In [17]: df.Property_Area.replace(['Urban', 'Rural', 'Semiurban'],[0,1,2],inplace=True)
```

```
In [18]: df.Loan_Status.replace(['Y', 'N'],[1,0],inplace=True)
```

```
In [19]: df.head()
```

```
Out[19]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantInc
0	0	0	0	0	0	5849	
1	0	1	1	0	0	4583	15
2	0	1	0	0	1	3000	
3	0	1	0	1	0	2583	23
4	0	0	0	0	0	6000	

```
In [20]: for col in cat_cols:
         df[col] = df[col].astype('int')
```

```
In [21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                614 non-null   int64
1   Married               614 non-null   int64
2   Dependents            614 non-null   int64
3   Education             614 non-null   int64
4   Self_Employed         614 non-null   int64
5   ApplicantIncome       614 non-null   int64
6   CoapplicantIncome     614 non-null   float64
7   LoanAmount            614 non-null   float64
8   Loan_Amount_Term      614 non-null   float64
9   Credit_History        614 non-null   float64
10  Property_Area         614 non-null   int64
11  Loan_Status           614 non-null   int64
dtypes: float64(4), int64(8)
memory usage: 57.7 KB
```

```
In [22]: from imblearn.combine import SMOTETomek
```

```
smote = SMOTETomek()

x=df.drop(columns=['Loan_Status'],axis=1)
y=df['Loan_Status']

x_bal,y_bal = smote.fit_resample(x,y)

print(y.value_counts())
print(y_bal.value_counts())
```

```

1    422
0    192
Name: Loan_Status, dtype: int64
1    358
0    358
Name: Loan_Status, dtype: int64

```

EDA - EXPLORATORY DATA ANALYSIS :-

```
In [23]: df.describe()
```


```
Out[23]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Credit_History
count	614.000000	614.000000	614.000000	614.000000	614.000000	614.000000	614.000000
mean	0.182410	0.653094	0.744300	0.218241	0.133550	5403.459283	0.960912
std	0.386497	0.476373	1.009623	0.413389	0.340446	6109.041673	0.200000
min	0.000000	0.000000	0.000000	0.000000	0.000000	150.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	2877.500000	0.000000
50%	0.000000	1.000000	0.000000	0.000000	0.000000	3812.500000	0.000000
75%	0.000000	1.000000	1.000000	0.000000	0.000000	5795.000000	0.000000
max	1.000000	1.000000	3.000000	1.000000	1.000000	81000.000000	1.000000

UNIVARIATE ANALYSIS :-


```
In [24]: sns.distplot(df['ApplicantIncome'], color='r')
```

```
Out[24]: <Axes: xlabel='ApplicantIncome', ylabel='Density'>
```





```
In [25]: sns.distplot(df['Credit_History'])
```

Out[25]: <Axes: xlabel='Credit_History', ylabel='Density'>




```
In [26]: sns.countplot(df,x='Gender')
plt.xlabel('Gender')
plt.show()
sns.countplot(df,x='Education')
plt.xlabel('Education')
plt.show()
```



```
In [27]: sns.countplot(df,x='Married')
```


```
Out[27]: <Axes: xlabel='Married', ylabel='count'>
```



BIVARIATE ANALYSIS :-


```
In [28]: plt.figure(figsize=(20,5))
plt.subplot(1,3,1)
sns.countplot(x= 'Married', hue = "Gender", data = df)
plt.subplot(1,3,2)
sns.countplot(x = 'Self_Employed', hue = "Education", data = df)
plt.subplot(1,3,3)
sns.countplot(x= 'Property_Area', hue = "Loan_Amount_Term", data =df)
```

Out[28]: <Axes: xlabel='Property_Area', ylabel='count'>



```
In [29]: sns.swarmplot(x='Gender', y='ApplicantIncome', hue = "Loan_Status", data
```

Out[29]: <Axes: xlabel='Gender', ylabel='ApplicantIncome'>



SCALING THE DATA :-

```
In [30]: scaler = StandardScaler()
x_bal = scaler.fit_transform(x_bal)
x_bal = pd.DataFrame(x_bal, columns=scaler.get_feature_names_out())
x_bal.head()
```

```
Out[30]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	-0.441956	-1.137924	-0.705071	-0.475423	-0.334367	0.152709	-
1	-0.441956	0.878793	0.357731	-0.475423	-0.334367	-0.090557	-
2	-0.441956	0.878793	-0.705071	-0.475423	2.990726	-0.394735	-
3	-0.441956	0.878793	-0.705071	2.103388	-0.334367	-0.474863	-
4	-0.441956	-1.137924	-0.705071	-0.475423	-0.334367	0.181724	-

SPLITTING THE DATA INTO TRAINING AND TESTING DATA :-

```
In [31]: x_train,x_test,y_train,y_test=train_test_split(x_bal,y_bal, test_size=0.3)
print(x_train.shape,x_test.shape)

(479, 11) (237, 11)
```

MODEL BUILDING :-

```
In [32]: def fit_model(model,name):
model.fit(x_train,y_train)

y_pred = model.predict(x_train)
print('training accuracy of ',name,' : ',accuracy_score(y_pred,y_train))
```

```

y_pred = model.predict(x_test)
print('testing accuracy of ',name,' : ',accuracy_score(y_pred, y_test))

plt.figure(figsize=(4,2))
sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
plt.show()

print(classification_report(y_test,y_pred))

```

```


In [33]: dtc = DecisionTreeClassifier()
fit_model(dtc,"DTC")

```

```

training accuracy of DTC : 1.0
testing accuracy of DTC : 0.7679324894514767

```



	precision	recall	f1-score	support
0	0.74	0.81	0.77	114
1	0.80	0.73	0.77	123
accuracy			0.77	237
macro avg	0.77	0.77	0.77	237
weighted avg	0.77	0.77	0.77	237

```


In [34]: rfc = RandomForestClassifier()
fit_model(rfc,"RFC")

```

```

training accuracy of RFC : 1.0
testing accuracy of RFC : 0.8523206751054853


```



	precision	recall	f1-score	support
0	0.86	0.82	0.84	114
1	0.84	0.88	0.86	123
accuracy			0.85	237
macro avg	0.85	0.85	0.85	237
weighted avg	0.85	0.85	0.85	237

```
In [35]: knn = KNeighborsClassifier()
fit_model(knn, 'KNN')
```


training accuracy of KNN : 0.824634655532359
testing accuracy of KNN : 0.7468354430379747



	precision	recall	f1-score	support
0	0.77	0.68	0.72	114
1	0.73	0.81	0.77	123
accuracy			0.75	237
macro avg	0.75	0.74	0.74	237
weighted avg	0.75	0.75	0.75	237

```
In [36]: from xgboost import XGBClassifier
xgb = XGBClassifier()
fit_model(xgb, 'XGB')
```

training accuracy of XGB : 1.0
testing accuracy of XGB : 0.8734177215189873



	precision	recall	f1-score	support
0	0.87	0.87	0.87	114
1	0.88	0.88	0.88	123
accuracy			0.87	237
macro avg	0.87	0.87	0.87	237
weighted avg	0.87	0.87	0.87	237

```
In [37]: ann = Sequential()
ann.add(Dense(units=12,activation='relu'))
ann.add(Dense(units=24,activation='relu'))
ann.add(Dense(units=1,activation='sigmoid'))
ann.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
ann.fit(x_train,y_train,batch_size=100,validation_split=0.2,epochs=100)
```

Epoch 1/100
4/4 [=====] - 1s 75ms/step - loss: 0.7652 - accuracy: 0.3708 - val_loss: 0.7232 - val_accuracy: 0.4271
Epoch 2/100
4/4 [=====] - 0s 11ms/step - loss: 0.7517 - accuracy: 0.3812 - val_loss: 0.7183 - val_accuracy: 0.4479
Epoch 3/100
4/4 [=====] - 0s 13ms/step - loss: 0.7388 - accuracy: 0.4178 - val_loss: 0.7138 - val_accuracy: 0.4167
Epoch 4/100
4/4 [=====] - 0s 10ms/step - loss: 0.7280 - accuracy: 0.4439 - val_loss: 0.7099 - val_accuracy: 0.4375
Epoch 5/100
4/4 [=====] - 0s 10ms/step - loss: 0.7166 - accuracy: 0.4909 - val_loss: 0.7057 - val_accuracy: 0.4688
Epoch 6/100
4/4 [=====] - 0s 11ms/step - loss: 0.7073 - accuracy: 0.5091 - val_loss: 0.7019 - val_accuracy: 0.4688
Epoch 7/100
4/4 [=====] - 0s 15ms/step - loss: 0.6985 - accuracy: 0.5274 - val_loss: 0.6981 - val_accuracy: 0.5208
Epoch 8/100
4/4 [=====] - 0s 11ms/step - loss: 0.6905 - accuracy: 0.5352 - val_loss: 0.6946 - val_accuracy: 0.5729
Epoch 9/100
4/4 [=====] - 0s 10ms/step - loss: 0.6823 - accuracy: 0.5587 - val_loss: 0.6910 - val_accuracy: 0.5521
Epoch 10/100
4/4 [=====] - 0s 10ms/step - loss: 0.6741 - accuracy: 0.5979 - val_loss: 0.6872 - val_accuracy: 0.5729
Epoch 11/100
4/4 [=====] - 0s 10ms/step - loss: 0.6664 - accuracy: 0.6162 - val_loss: 0.6836 - val_accuracy: 0.6042
Epoch 12/100
4/4 [=====] - 0s 11ms/step - loss: 0.6593 - accuracy: 0.6397 - val_loss: 0.6800 - val_accuracy: 0.6146
Epoch 13/100
4/4 [=====] - 0s 14ms/step - loss: 0.6516 - accuracy: 0.6423 - val_loss: 0.6763 - val_accuracy: 0.6250
Epoch 14/100
4/4 [=====] - 0s 11ms/step - loss: 0.6445 - accuracy: 0.6632 - val_loss: 0.6727 - val_accuracy: 0.6458
Epoch 15/100
4/4 [=====] - 0s 10ms/step - loss: 0.6371 - accuracy: 0.6684 - val_loss: 0.6689 - val_accuracy: 0.6458
Epoch 16/100
4/4 [=====] - 0s 10ms/step - loss: 0.6296 - accuracy: 0.6971 - val_loss: 0.6650 - val_accuracy: 0.6354
Epoch 17/100
4/4 [=====] - 0s 10ms/step - loss: 0.6226 - accuracy: 0.6997 - val_loss: 0.6613 - val_accuracy: 0.6458
Epoch 18/100
4/4 [=====] - 0s 10ms/step - loss: 0.6154 - accuracy: 0.7076 - val_loss: 0.6578 - val_accuracy: 0.6458
Epoch 19/100
4/4 [=====] - 0s 11ms/step - loss: 0.6086 - accuracy: 0.7232 - val_loss: 0.6548 - val_accuracy: 0.6458
Epoch 20/100
4/4 [=====] - 0s 11ms/step - loss: 0.6018 - accuracy: 0.7232 - val_loss: 0.6512 - val_accuracy: 0.6562

Epoch 21/100
4/4 [=====] - 0s 17ms/step - loss: 0.5947 - accuracy: 0.7311 - val_loss: 0.6481 - val_accuracy: 0.6562
Epoch 22/100
4/4 [=====] - 0s 11ms/step - loss: 0.5877 - accuracy: 0.7363 - val_loss: 0.6449 - val_accuracy: 0.6562
Epoch 23/100
4/4 [=====] - 0s 11ms/step - loss: 0.5813 - accuracy: 0.7493 - val_loss: 0.6422 - val_accuracy: 0.6458
Epoch 24/100
4/4 [=====] - 0s 11ms/step - loss: 0.5743 - accuracy: 0.7598 - val_loss: 0.6394 - val_accuracy: 0.6458
Epoch 25/100
4/4 [=====] - 0s 17ms/step - loss: 0.5675 - accuracy: 0.7650 - val_loss: 0.6368 - val_accuracy: 0.6562
Epoch 26/100
4/4 [=====] - 0s 11ms/step - loss: 0.5609 - accuracy: 0.7702 - val_loss: 0.6347 - val_accuracy: 0.6562
Epoch 27/100
4/4 [=====] - 0s 18ms/step - loss: 0.5545 - accuracy: 0.7702 - val_loss: 0.6330 - val_accuracy: 0.6562
Epoch 28/100
4/4 [=====] - 0s 11ms/step - loss: 0.5483 - accuracy: 0.7728 - val_loss: 0.6310 - val_accuracy: 0.6458
Epoch 29/100
4/4 [=====] - 0s 10ms/step - loss: 0.5421 - accuracy: 0.7755 - val_loss: 0.6292 - val_accuracy: 0.6458
Epoch 30/100
4/4 [=====] - 0s 10ms/step - loss: 0.5364 - accuracy: 0.7807 - val_loss: 0.6276 - val_accuracy: 0.6458
Epoch 31/100
4/4 [=====] - 0s 10ms/step - loss: 0.5305 - accuracy: 0.7807 - val_loss: 0.6265 - val_accuracy: 0.6458
Epoch 32/100
4/4 [=====] - 0s 10ms/step - loss: 0.5249 - accuracy: 0.7833 - val_loss: 0.6253 - val_accuracy: 0.6458
Epoch 33/100
4/4 [=====] - 0s 11ms/step - loss: 0.5195 - accuracy: 0.7833 - val_loss: 0.6246 - val_accuracy: 0.6458
Epoch 34/100
4/4 [=====] - 0s 13ms/step - loss: 0.5144 - accuracy: 0.7833 - val_loss: 0.6235 - val_accuracy: 0.6562
Epoch 35/100
4/4 [=====] - 0s 11ms/step - loss: 0.5098 - accuracy: 0.7833 - val_loss: 0.6232 - val_accuracy: 0.6562
Epoch 36/100
4/4 [=====] - 0s 10ms/step - loss: 0.5047 - accuracy: 0.7859 - val_loss: 0.6228 - val_accuracy: 0.6562
Epoch 37/100
4/4 [=====] - 0s 10ms/step - loss: 0.5002 - accuracy: 0.7833 - val_loss: 0.6226 - val_accuracy: 0.6458
Epoch 38/100
4/4 [=====] - 0s 10ms/step - loss: 0.4955 - accuracy: 0.7833 - val_loss: 0.6221 - val_accuracy: 0.6458
Epoch 39/100
4/4 [=====] - 0s 10ms/step - loss: 0.4916 - accuracy: 0.7885 - val_loss: 0.6215 - val_accuracy: 0.6458
Epoch 40/100
4/4 [=====] - 0s 15ms/step - loss: 0.4875 - accuracy: 0.7911 - val_loss: 0.6214 - val_accuracy: 0.6354

Epoch 41/100
4/4 [=====] - 0s 19ms/step - loss: 0.4838 - accuracy: 0.7911 - val_loss: 0.6216 - val_accuracy: 0.6354
Epoch 42/100
4/4 [=====] - 0s 12ms/step - loss: 0.4800 - accuracy: 0.7885 - val_loss: 0.6214 - val_accuracy: 0.6354
Epoch 43/100
4/4 [=====] - 0s 19ms/step - loss: 0.4767 - accuracy: 0.7911 - val_loss: 0.6210 - val_accuracy: 0.6458
Epoch 44/100
4/4 [=====] - 0s 11ms/step - loss: 0.4734 - accuracy: 0.7911 - val_loss: 0.6206 - val_accuracy: 0.6458
Epoch 45/100
4/4 [=====] - 0s 11ms/step - loss: 0.4705 - accuracy: 0.7937 - val_loss: 0.6213 - val_accuracy: 0.6458
Epoch 46/100
4/4 [=====] - 0s 11ms/step - loss: 0.4672 - accuracy: 0.7911 - val_loss: 0.6215 - val_accuracy: 0.6458
Epoch 47/100
4/4 [=====] - 0s 12ms/step - loss: 0.4648 - accuracy: 0.7911 - val_loss: 0.6217 - val_accuracy: 0.6458
Epoch 48/100
4/4 [=====] - 0s 11ms/step - loss: 0.4621 - accuracy: 0.7911 - val_loss: 0.6223 - val_accuracy: 0.6458
Epoch 49/100
4/4 [=====] - 0s 13ms/step - loss: 0.4597 - accuracy: 0.7911 - val_loss: 0.6227 - val_accuracy: 0.6458
Epoch 50/100
4/4 [=====] - 0s 11ms/step - loss: 0.4573 - accuracy: 0.7963 - val_loss: 0.6234 - val_accuracy: 0.6458
Epoch 51/100
4/4 [=====] - 0s 16ms/step - loss: 0.4552 - accuracy: 0.7990 - val_loss: 0.6230 - val_accuracy: 0.6458
Epoch 52/100
4/4 [=====] - 0s 11ms/step - loss: 0.4530 - accuracy: 0.8042 - val_loss: 0.6232 - val_accuracy: 0.6458
Epoch 53/100
4/4 [=====] - 0s 10ms/step - loss: 0.4508 - accuracy: 0.8016 - val_loss: 0.6236 - val_accuracy: 0.6458
Epoch 54/100
4/4 [=====] - 0s 11ms/step - loss: 0.4491 - accuracy: 0.8042 - val_loss: 0.6239 - val_accuracy: 0.6458
Epoch 55/100
4/4 [=====] - 0s 11ms/step - loss: 0.4471 - accuracy: 0.8016 - val_loss: 0.6242 - val_accuracy: 0.6458
Epoch 56/100
4/4 [=====] - 0s 11ms/step - loss: 0.4455 - accuracy: 0.8016 - val_loss: 0.6244 - val_accuracy: 0.6458
Epoch 57/100
4/4 [=====] - 0s 11ms/step - loss: 0.4437 - accuracy: 0.8016 - val_loss: 0.6240 - val_accuracy: 0.6562
Epoch 58/100
4/4 [=====] - 0s 11ms/step - loss: 0.4419 - accuracy: 0.8042 - val_loss: 0.6242 - val_accuracy: 0.6667
Epoch 59/100
4/4 [=====] - 0s 11ms/step - loss: 0.4405 - accuracy: 0.8042 - val_loss: 0.6253 - val_accuracy: 0.6667
Epoch 60/100
4/4 [=====] - 0s 11ms/step - loss: 0.4391 - accuracy: 0.8068 - val_loss: 0.6256 - val_accuracy: 0.6667

Epoch 61/100
4/4 [=====] - 0s 11ms/step - loss: 0.4378 - accuracy: 0.8068 - val_loss: 0.6252 - val_accuracy: 0.6667
Epoch 62/100
4/4 [=====] - 0s 11ms/step - loss: 0.4364 - accuracy: 0.8068 - val_loss: 0.6261 - val_accuracy: 0.6667
Epoch 63/100
4/4 [=====] - 0s 10ms/step - loss: 0.4352 - accuracy: 0.8068 - val_loss: 0.6274 - val_accuracy: 0.6667
Epoch 64/100
4/4 [=====] - 0s 10ms/step - loss: 0.4338 - accuracy: 0.8068 - val_loss: 0.6269 - val_accuracy: 0.6667
Epoch 65/100
4/4 [=====] - 0s 10ms/step - loss: 0.4326 - accuracy: 0.8068 - val_loss: 0.6271 - val_accuracy: 0.6667
Epoch 66/100
4/4 [=====] - 0s 11ms/step - loss: 0.4316 - accuracy: 0.8042 - val_loss: 0.6279 - val_accuracy: 0.6667
Epoch 67/100
4/4 [=====] - 0s 16ms/step - loss: 0.4305 - accuracy: 0.8068 - val_loss: 0.6280 - val_accuracy: 0.6667
Epoch 68/100
4/4 [=====] - 0s 11ms/step - loss: 0.4296 - accuracy: 0.8042 - val_loss: 0.6285 - val_accuracy: 0.6667
Epoch 69/100
4/4 [=====] - 0s 12ms/step - loss: 0.4287 - accuracy: 0.8042 - val_loss: 0.6287 - val_accuracy: 0.6667
Epoch 70/100
4/4 [=====] - 0s 11ms/step - loss: 0.4275 - accuracy: 0.8042 - val_loss: 0.6278 - val_accuracy: 0.6667
Epoch 71/100
4/4 [=====] - 0s 10ms/step - loss: 0.4266 - accuracy: 0.8068 - val_loss: 0.6287 - val_accuracy: 0.6667
Epoch 72/100
4/4 [=====] - 0s 13ms/step - loss: 0.4256 - accuracy: 0.8068 - val_loss: 0.6289 - val_accuracy: 0.6667
Epoch 73/100
4/4 [=====] - 0s 11ms/step - loss: 0.4249 - accuracy: 0.8068 - val_loss: 0.6289 - val_accuracy: 0.6667
Epoch 74/100
4/4 [=====] - 0s 11ms/step - loss: 0.4239 - accuracy: 0.8068 - val_loss: 0.6308 - val_accuracy: 0.6667
Epoch 75/100
4/4 [=====] - 0s 12ms/step - loss: 0.4233 - accuracy: 0.8042 - val_loss: 0.6331 - val_accuracy: 0.6667
Epoch 76/100
4/4 [=====] - 0s 10ms/step - loss: 0.4221 - accuracy: 0.8042 - val_loss: 0.6327 - val_accuracy: 0.6667
Epoch 77/100
4/4 [=====] - 0s 11ms/step - loss: 0.4213 - accuracy: 0.8042 - val_loss: 0.6322 - val_accuracy: 0.6667
Epoch 78/100
4/4 [=====] - 0s 11ms/step - loss: 0.4206 - accuracy: 0.8042 - val_loss: 0.6313 - val_accuracy: 0.6667
Epoch 79/100
4/4 [=====] - 0s 10ms/step - loss: 0.4195 - accuracy: 0.8042 - val_loss: 0.6323 - val_accuracy: 0.6667
Epoch 80/100
4/4 [=====] - 0s 11ms/step - loss: 0.4188 - accuracy: 0.8042 - val_loss: 0.6329 - val_accuracy: 0.6667

Epoch 81/100
4/4 [=====] - 0s 11ms/step - loss: 0.4178 - accuracy: 0.8042 - val_loss: 0.6330 - val_accuracy: 0.6667
Epoch 82/100
4/4 [=====] - 0s 10ms/step - loss: 0.4170 - accuracy: 0.8042 - val_loss: 0.6333 - val_accuracy: 0.6667
Epoch 83/100
4/4 [=====] - 0s 11ms/step - loss: 0.4163 - accuracy: 0.8042 - val_loss: 0.6339 - val_accuracy: 0.6667
Epoch 84/100
4/4 [=====] - 0s 11ms/step - loss: 0.4155 - accuracy: 0.8042 - val_loss: 0.6336 - val_accuracy: 0.6667
Epoch 85/100
4/4 [=====] - 0s 13ms/step - loss: 0.4148 - accuracy: 0.8042 - val_loss: 0.6335 - val_accuracy: 0.6667
Epoch 86/100
4/4 [=====] - 0s 14ms/step - loss: 0.4141 - accuracy: 0.8042 - val_loss: 0.6341 - val_accuracy: 0.6667
Epoch 87/100
4/4 [=====] - 0s 12ms/step - loss: 0.4136 - accuracy: 0.8068 - val_loss: 0.6345 - val_accuracy: 0.6667
Epoch 88/100
4/4 [=====] - 0s 11ms/step - loss: 0.4128 - accuracy: 0.8068 - val_loss: 0.6344 - val_accuracy: 0.6667
Epoch 89/100
4/4 [=====] - 0s 11ms/step - loss: 0.4122 - accuracy: 0.8068 - val_loss: 0.6341 - val_accuracy: 0.6667
Epoch 90/100
4/4 [=====] - 0s 11ms/step - loss: 0.4116 - accuracy: 0.8068 - val_loss: 0.6348 - val_accuracy: 0.6667
Epoch 91/100
4/4 [=====] - 0s 11ms/step - loss: 0.4109 - accuracy: 0.8068 - val_loss: 0.6348 - val_accuracy: 0.6667
Epoch 92/100
4/4 [=====] - 0s 11ms/step - loss: 0.4103 - accuracy: 0.8068 - val_loss: 0.6358 - val_accuracy: 0.6667
Epoch 93/100
4/4 [=====] - 0s 11ms/step - loss: 0.4097 - accuracy: 0.8068 - val_loss: 0.6357 - val_accuracy: 0.6667
Epoch 94/100
4/4 [=====] - 0s 11ms/step - loss: 0.4090 - accuracy: 0.8068 - val_loss: 0.6374 - val_accuracy: 0.6667
Epoch 95/100
4/4 [=====] - 0s 15ms/step - loss: 0.4085 - accuracy: 0.8094 - val_loss: 0.6385 - val_accuracy: 0.6667
Epoch 96/100
4/4 [=====] - 0s 11ms/step - loss: 0.4079 - accuracy: 0.8094 - val_loss: 0.6374 - val_accuracy: 0.6667
Epoch 97/100
4/4 [=====] - 0s 11ms/step - loss: 0.4074 - accuracy: 0.8094 - val_loss: 0.6382 - val_accuracy: 0.6667
Epoch 98/100
4/4 [=====] - 0s 12ms/step - loss: 0.4067 - accuracy: 0.8094 - val_loss: 0.6371 - val_accuracy: 0.6667
Epoch 99/100
4/4 [=====] - 0s 16ms/step - loss: 0.4060 - accuracy: 0.8120 - val_loss: 0.6376 - val_accuracy: 0.6771
Epoch 100/100
4/4 [=====] - 0s 10ms/step - loss: 0.4055 - accuracy: 0.8120 - val_loss: 0.6394 - val_accuracy: 0.6771

Out[37]: <keras.callbacks.History at 0x7f8615ca4e20>

```
In [38]: y_pred = ann.predict(x_train)
y_pred = y_pred > 0.5
print('training accuracy of ANN : ',accuracy_score(y_pred,y_train))

y_pred = ann.predict(x_test)
y_pred = y_pred > 0.5

print('testing accuracy of ANN : ',accuracy_score(y_pred, y_test))

15/15 [=====] - 0s 1ms/step
training accuracy of ANN : 0.7828810020876826
8/8 [=====] - 0s 2ms/step
testing accuracy of ANN : 0.7848101265822784
```


HYPER PARAMETER TUNING:-

```
In [39]: from sklearn.model_selection import GridSearchCV
params = {
    'criterion' :['gini','entropy'],
    'max_depth' :[None,5,10,15],
    'min_samples_split':[2,5,10],
    'min_samples_leaf':[1,2,4]
}
gcv = GridSearchCV(DecisionTreeClassifier(),params,cv=5)
gcv.fit(x_train,y_train)
gcv.best_params_
```

Out[39]: {'criterion': 'entropy',
 'max_depth': 5,
 'min_samples_leaf': 4,
 'min_samples_split': 10}

```
In [51]: dtc2 = DecisionTreeClassifier(criterion='entropy',max_depth=5,min_samples
fit_model(dtc2,'DTC after tuning')
```

training accuracy of DTC after tuning : 0.8121085594989561
testing accuracy of DTC after tuning : 0.7637130801687764




	precision	recall	f1-score	support
0	0.84	0.63	0.72	114
1	0.72	0.89	0.80	123
accuracy			0.76	237
macro avg	0.78	0.76	0.76	237
weighted avg	0.78	0.76	0.76	237

```
In [41]: parameters = {
        'n_estimators' : [1,20,30,55,68,74,90,120,115],
        'criterion' : ['gini','entropy'],
        'max_features' : [ "sqrt" , "log2"],
        'max_depth' : [2,5,8,10]
    }
    rcv = RandomizedSearchCV(estimator=RandomForestClassifier(),param_distrib
    rcv.fit(x_train,y_train)
    rcv.best_params_
```

```
Out[41]: {'n_estimators': 90,
        'max_features': 'log2',
        'max_depth': 8,
        'criterion': 'gini'}
```

```
In [50]: rfc2 = RandomForestClassifier(n_estimators=90,max_features='log2',max_dep
    fit_model(rfc2,'RFC ( after tuning ) ')
```

```
training accuracy of  RFC ( after tuning )   : 0.9373695198329853
testing accuracy of  RFC ( after tuning )   : 0.8270042194092827
```



	precision	recall	f1-score	support
0	0.87	0.75	0.81	114
1	0.80	0.89	0.84	123
accuracy			0.83	237
macro avg	0.83	0.82	0.83	237
weighted avg	0.83	0.83	0.83	237

```
In [43]: param_grid = {
        'n_neighbors': [3, 5, 7],
        'weights': ['uniform', 'distance'],
        'p': [1, 2]
    }


    g2 = GridSearchCV(estimator=KNeighborsClassifier(), param_grid=param_grid
    g2.fit(x_train, y_train)

    print( g2.best_params_ )

    {'n_neighbors': 5, 'p': 1, 'weights': 'distance'}
```

```
In [49]: knn2 = KNeighborsClassifier(n_neighbors=5,p=1,weights='distance')
    fit_model(knn2,'KNN (after tuning)')
```

```
training accuracy of  KNN (after tuning)   : 1.0
testing accuracy of  KNN (after tuning)   : 0.810126582278481
```



	precision	recall	f1-score	support
0	0.85	0.74	0.79	114
1	0.78	0.88	0.83	123
accuracy			0.81	237
macro avg	0.82	0.81	0.81	237
weighted avg	0.81	0.81	0.81	237


```
In [45]: param_grid = {
          'learning_rate': [0.1, 0.2, 0.3],
          'max_depth': [3, 4, 5],
          'n_estimators': [100, 200, 300]
        }
g = GridSearchCV(estimator=XGBClassifier(), param_grid=param_grid, cv=5,
g.fit(x_train, y_train)

g.best_params_
```

```
Out[45]: {'learning_rate': 0.2, 'max_depth': 5, 'n_estimators': 100}
```

```
In [48]: xgb2 = XGBClassifier(learning_rate = 0.2,max_depth=5,n_estimators = 100)
fit_model(xgb2,'XGB (after tuning)')
```

```
training accuracy of XGB (after tuning) : 1.0
testing accuracy of XGB (after tuning) : 0.8649789029535865
```



	precision	recall	f1-score	support
0	0.87	0.84	0.86	114
1	0.86	0.89	0.87	123
accuracy			0.86	237
macro avg	0.87	0.86	0.86	237
weighted avg	0.87	0.86	0.86	237

SAVING THE MODEL :-

```
In [47]: pickle.dump(xgb2,open('xgb.pkl','wb'))
pickle.dump(scaler,open('scaler.pkl','wb'))
```